# Environmental fluctuation strengthens gene regulation only when it can be tracked

[**PENDING:** Author names][1*]

**1** [**PENDING:** Affiliation]

* [**PENDING:** corresponding author e-mail]

## Abstract

Environmental variability is widely expected to favour more elaborate gene regulation, but models disagree even about the direction of the effect, and many of them measure regulatory complexity by counting network connections. We separated fluctuation from the information available for tracking it and asked how each affects regulatory architecture, using two independent individual-based models of genome evolution. In EvoSim the genome is a regulatory network of evolvable size, and in R-Aevol regulation arises from transcription factors encoded in a DNA-like sequence. In each model we compared frozen and fluctuating environments and, under fluctuation, an informative cue, no cue, and a yoked cue that supplies the same signals at the same frequencies but independently of the environment. In EvoSim, fluctuation increased the density of regulatory connections by 26% but did not change total wiring strength, so the average connection became weaker. In R-Aevol we tested four predictions, committed to version control before the simulations were run, at three switching rates, and all four were supported. With an informative cue, fluctuation increased total regulatory strength 1.8- to 2.7-fold relative to a frozen environment, and informative cues produced 2.3 to 2.9 times the regulatory strength of yoked cues but no more genes. With an informative cue, fluctuation also reduced the number of genes by 8 to 10. Without a cue, fluctuating populations had fewer genes and the weakest regulation of any condition, whether an environment lasted half a lifetime or about 250 generations. In neither model did connection counts track regulatory strength. With an informative cue, R-Aevol populations had no more regulatory links than frozen populations, although their regulation was stronger. Fluctuation therefore strengthened regulation only when the environment could be tracked, and counting edges can miss the effect.

## Author summary

Organisms in changing environments often have complex gene regulation, and it is easy to assume that environmental change is what builds it. We checked this assumption in two computer models of evolving genomes that work in very different ways. Populations evolved in environments that were either constant or switching between states, and we controlled what the organisms could know about the switching. Some populations received a signal that told them which state they were in. Others received the same signals arranged so that they told them nothing, and a third group received no signal. Switching on its own did not produce stronger regulation. In the model in which organisms used the signal, a meaningful signal made regulatory connections much stronger than in a constant environment. Without a signal, they became weaker than in any other condition, whether the environment switched within

a lifetime or only about every 250 generations. The uninformative signal did not help, although it was physically identical to the meaningful one. We also found that counting regulatory connections, a common measure of network complexity, can miss changes in how strong the connections are. Our main predictions were written down and committed before the simulations were run.

# Introduction

Organisms that live in variable environments are often assumed to need more elaborate gene regulation, and therefore to evolve it. The argument is simple. Regulation lets one genome produce different phenotypes in different conditions, and in a constant environment there is little for it to do. Comparative genomics offers some support. In bacteria the number of transcription factors grows faster than the number of genes [1], roughly as its square [2]. Free-living species carry more transcription factors than their genome size predicts [3], and intracellular parasites and other species of stable niches have simpler signalling systems than species from the open environment [4]. Simulations with R-Aevol, however, reproduce such scaling laws, trace them to differences in mutation and rearrangement rates between lineages and find that lifestyle does not explain them [5], so the comparative evidence alone does not show that variable environments select for regulation.

At the level of the phenotype, theory specifies when variability favours a response. A plastic response within a lifetime pays when the environment changes on the timescale of an individual's life and a reliable cue indicates the coming state [6–9]. Without such a cue, fluctuation favours a fixed generalist phenotype, bet-hedging or genetic tracking [8,9], and responding to a cue that predicts the environment poorly amplifies the mismatch between phenotype and environment [10]. A recent model of a gene regulatory network shows the same pattern for the evolution of plasticity. Plasticity evolved mainly under fast, erratic change, especially when cues were reliable, and evolvability under intermediate variability with less reliable cues [11].

Much less is known about the consequences for regulatory architecture, by which we mean the number of genes and the strength with which they regulate one another, and on this point models disagree. Tsuda and Kawata found that fluctuating selection promoted complex regulatory networks, mostly through the fixation of beneficial gene duplications, when the optimum moved in a random walk. When it alternated between two positions the effect was slight and limited to large, infrequent shifts [12]. Their networks received an external signal that told them when to express the phenotype but carried no information about where the optimum was. Bentkowski and colleagues reported larger genomes in fluctuating environments [13]. In a model in which network size was fixed, however, smaller networks let populations persist longer when the environment changed fast [14]. Studies with the Aevol platform, on the other hand, find genome size limited by selection for mutational robustness and reduced by large population size and high mutation rates [15,16]. Work on modularity shows that its evolution depends on sparseness and on how the variation is structured [17,18].

Differences of method make these results hard to compare. Fluctuation and information are usually confounded. An environment that changes more often also contains a different amount of predictable structure, and adding a cue changes what an organism experiences as well as what it can know. A comparison of networks with and without a cue input [19,20] changes the two together. Studies of animal learning handle this problem with a yoked control, in which a second subject receives the same stimuli as the first but without the contingency between stimulus and outcome. Church showed that pairing individual subjects in this way can bias the comparison [21]. In our design each yoked arm receives the same stimulus statistics as the informative arm without being matched to particular informative-arm populations, so this bias does not arise. Varying the reliability of a cue achieves the same separation only if it includes a cue that is present at the usual frequencies but carries no information, and we know of no explicit use of a yoked control in studies of regulatory network evolution. Measurement is a further problem. Regulatory complexity is often quantified by counting genes or connections [12,17], but network size and connectivity can grow without selection for them [22], and a count does not reflect how strongly a connection acts. Weighted measures have been used [23], though seldom as the main outcome.

We addressed these issues with two independent models of genome evolution that represent genomes and regulation in different ways. In EvoSim the genome is a regulatory network whose size and connectivity evolve. In R-Aevol, regulation arises from transcription factors encoded in a DNA-like sequence that bind upstream of genes [5,24]. In each model we compared frozen and fluctuating environments, and under fluctuation we compared an informative cue, no cue, and a yoked cue that presents the same signals at the same frequencies as the informative one but independently of the environment (in R-Aevol also with the same per-age rates of target and signal change). Connectivity was

measured as a count and as a strength. An exploratory replication in R-Aevol contradicted one conclusion we had drawn earlier from EvoSim, that fluctuation enlarges genomes. It reproduced another, that a cue's information does not change genome size, and showed in addition that information changes regulation. We then stated four predictions, committed them before running new simulations, and tested them at three switching rates and two population sizes. A further pre-registered experiment varied environmental grain, one of the differences between our design and that of Tsuda and Kawata.

Fluctuation strengthened regulation when a cue made the environment trackable and weakened it when no cue was available. The difference came from the information the cue carried, since a cue that carried none had no such effect. Information did not increase genome size, and fluctuation reduced it. Without a cue, these effects did not detectably depend on whether the environment changed within a lifetime or about every 250 generations, so environmental grain does not explain the difference from Tsuda and Kawata's results. Connection counts did not track connection strength, so a count alone can misrepresent the effect of fluctuation on a network.

# Results

## Models and experimental design

**EvoSim** is an individual-based model of unicellular organisms on a two-dimensional lattice with diffusing, cross-fed resources (Materials and methods; ODD description in S1 Text). Each genome is a gene regulatory network, specified by a weight matrix and node biases, with 13 to 32 nodes. Thirteen nodes are fixed. Twelve are read out as phenotypic traits (resource-uptake allocation, body size, motility, division threshold, DNA repair, stress tolerance and leakage) and one receives an abiotic cue, much like the sensor gene of Sakamoto and Innan [25]. The remaining hidden nodes are added by duplication and removed by deletion, so genome size can evolve. The network is developed once at birth, and the phenotype does not change afterwards. Mutation can create and remove connections and each non-zero connection has a small maintenance cost, so connectivity is also free to evolve.

**R-Aevol** is the version of the Aevol platform that includes gene regulation [5, 24, 26]. Genomes are circular sequences in which promoters, genes and binding sites are identified by motif search. Translated proteins contribute to a phenotype, a function that is compared with a target function, and can also bind upstream of promoters to activate or repress the transcription of other genes. Protein concentrations are integrated over the 20 ages of an individual's life, and fitness is evaluated against the target at each age. Environmental signals are signal proteins that are present at a fixed concentration during the ages in which their environment is active, so organisms can evolve to respond to them.

The design (Fig 1) crosses environment (frozen or fluctuating) with cue (informative, absent or yoked). In R-Aevol the fluctuating environment switches between two target functions, A and B. B is the mirror image of A and has the same geometric area, so neither target is harder to match. In the frozen environment the switching mechanism is kept but the two targets are identical. In the informative arm each target is paired with its own signal protein. The yoked arm has four environments, one for each combination of target and signal, so the two signals occur equally often but say nothing about the target. At each switch R-Aevol moves to one of the other environments at random. At a given switching probability the four-environment arm thus changes its target less often than the two-environment arm, and we multiplied the yoked arm's switching probability by 1.5 so that the two arms have the same per-age rates of target change and of signal change (Materials and methods). We write $q$ for the switching probability per age. EvoSim follows the same logic with an abiotic cue, whose yoked version is decoupled from the environmental state by running on an incommensurate period.

## In EvoSim, fluctuation increases connection number but not wiring strength

Earlier EvoSim experiments had found that fluctuation increases connection density (S4 Text). We repeated the comparison in 180 runs (six arms of 30 seeds, 150,000 time steps each) and measured the strength of connections as well as their number. The cue is zero in a frozen environment, so the three frozen arms are identical and the runs contain 120 distinct populations. Without a cue, fluctuation raised

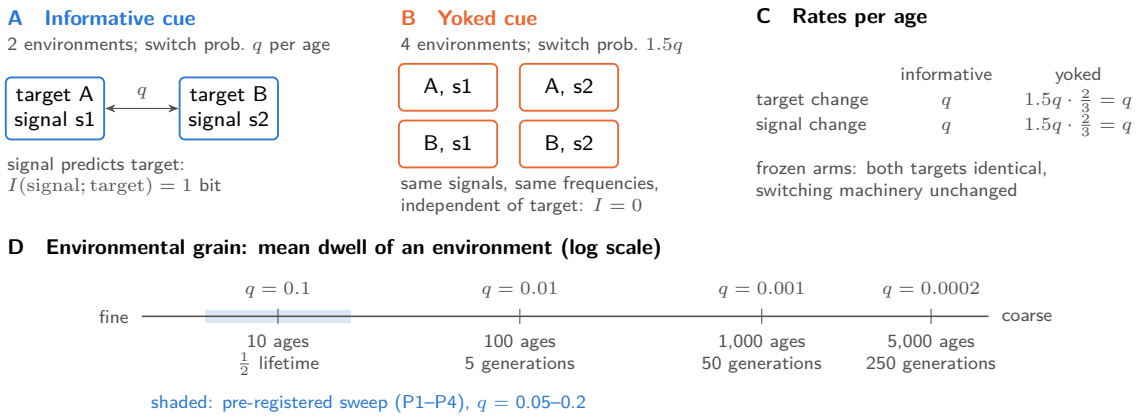



**Fig 1. Design.** (A) The informative arm in R-Aevol has two environments, each pairing a target with its own signal protein, so the signal predicts the target. (B) The yoked arm has four environments pairing each target with each signal, so the same signals occur at the same frequencies but carry no information about the target. (C) Because a switch moves to a different environment chosen at random, the yoked arm switches with probability $1.5q$, which gives it the same per-age rates of target change and signal change as the informative arm. Frozen arms use two identical targets. (D) Environmental grain as the mean dwell time of an environment, set by the switching probability $q$ over the 20 ages of a lifetime. Ticks mark the four switching probabilities of the grain experiment (G1–G4); shading marks the range of the pre-registered sweep (P1–P4).

count density, the proportion of possible connections that are non-zero, from 0.467 to 0.587 (+25.5%, $p < 10^{-4}$, $d = 2.2$; 28 surviving fluctuating and 30 frozen populations; Fig 2A), and the number of non-zero connections from 86 to 120 (+39%, $p < 10^{-4}$). Total wiring strength did not change (38.0 in frozen and 37.5 in fluctuating populations, $p = 0.86$; Fig 2C). Fluctuating networks thus had more connections but not more wiring, and the average non-zero connection was about 30% weaker. Strength density, the summed absolute weight divided by the number of possible connections, fell from 0.207 to 0.184 (−10.9%, $p < 10^{-4}$, $d = -1.4$; Fig 2B), mostly because there were more possible connections. Node count stayed close to its minimum of 13 in this regime and increased only slightly (+0.6 nodes, $p = 0.047$), which raises $n_g^2$ by about 9%. With the informative and the yoked cue, fluctuation had effects in the same directions (count density +18% and +28%, total strength unchanged, $p = 0.40$ and 0.59).

Population size is a possible confound. Fluctuating populations were 40% smaller than frozen ones, as expected since in frozen dishes all inflow feeds the best-paying channel (S1 Text), and population size can change genome architecture by itself. Because the population sizes of the two arms did not overlap, the two effects cannot be separated statistically, and the correlations within arms point in different directions for the two measures. Larger populations had lower strength density in every arm ($r$ between −0.31 and −0.47; Fig 2D), so population size would predict higher strength density in the smaller fluctuating populations, the reverse of what we found. Count density, in contrast, was higher in smaller populations within each fluctuating arm ($r$ between −0.63 and −0.72), though not in the frozen arm ($r = 0.24$), so population size may contribute to the increase in connection number. The cue was also not used in this regime. Populations with the informative and the yoked cue were the same size ($p = 0.79$) and had the same strength density ($p = 0.99$), although count density was lower with the informative cue (0.55 against 0.60, $p = 0.009$). EvoSim organisms respond to the cue only when development is reversible, and even then they match their phenotype to the current state only in part (S4 Text). The EvoSim results concern fluctuation that the organisms cannot track.

## In R-Aevol, the effect of fluctuation on regulation depends on the cue

We then asked whether an independent model would reproduce two earlier EvoSim conclusions, namely that fluctuation increases genome size and connectivity and that the information content of a cue makes no

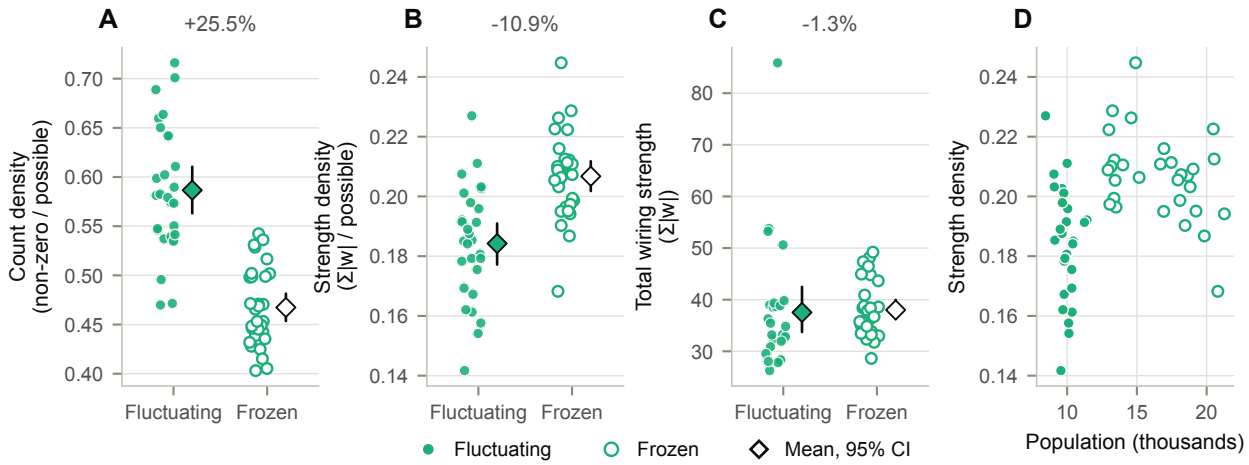



**Fig 2. EvoSim: more connections, no more wiring.** Fluctuating (no cue; filled) versus frozen (hollow) populations, one point per population; diamonds are means with bootstrap 95% confidence intervals. Numbers above panels are the change in the mean. (A) Count density, the number of non-zero connections divided by the number of possible connections. (B) Strength density, the summed absolute weight divided by the number of possible connections. (C) Total wiring strength. (D) Strength density against population size within each arm. Extinct populations (2 of 30 fluctuating, 0 of 30 frozen) are excluded.

difference to genome size. It reproduced the second but not the first. In an exploratory experiment (six arms, 12 seeds each, 30,000 generations), fluctuation reduced genome size under every cue condition, and its effect on regulation depended on the cue (Fig 3). With an informative cue, total regulatory strength was 12.7 in fluctuating populations and 4.7 in frozen populations given the same signals. Fluctuating populations without a cue had the lowest regulatory strength of the six arms (1.4).

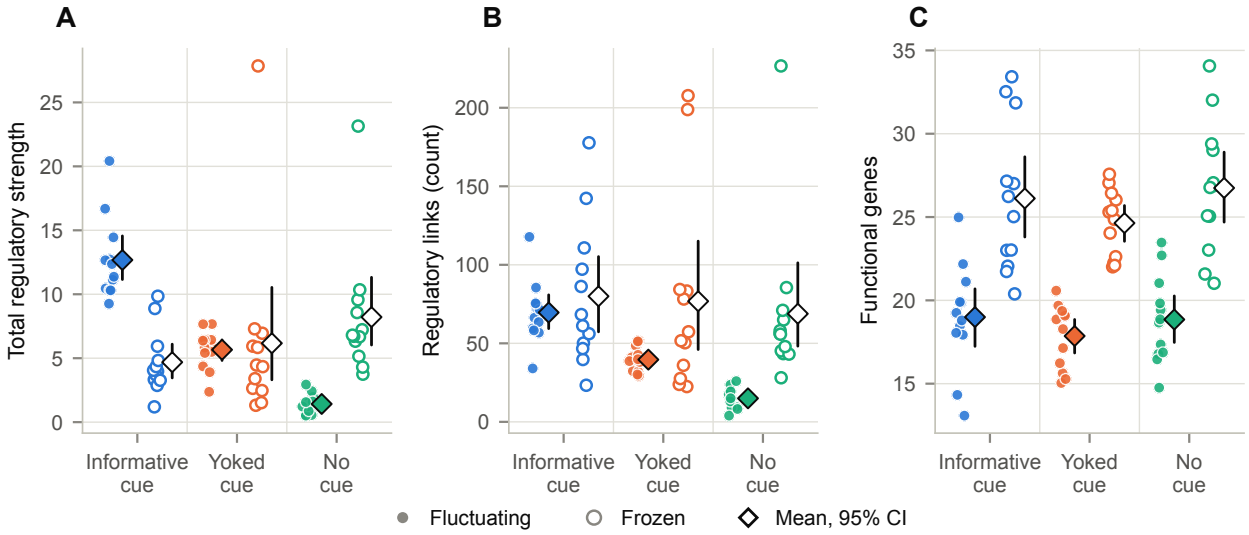



**Fig 3. R-Aevol, six arms (exploratory).** Total regulatory strength (A), number of regulatory links (B) and number of functional genes (C) for each combination of environment (filled, fluctuating; hollow, frozen) and cue (colour), one point per population; means with bootstrap 95% confidence intervals. $q = 0.10$, population 1024, 12 populations per arm.

Since this interpretation came from inspecting the data, we tested it with new seeds. Four predictions, with their measures, statistical tests and decision rules, were committed to version control before any of the

new simulations were run (Table 2; Materials and methods). The switching probability took three values over a fourfold range ($q = 0.05$, 0.10 and 0.20) at a population size of 1024, with 12 new seeds per condition, and the main contrasts were repeated at a population size of 512. For the two predictions that depend on cue use we also specified a manipulation check. At a given rate, the informative cue had to reduce metabolic error relative to the yoked cue, showing that the organisms used it. The check was met at $q = 0.10$ and $q = 0.20$ ($p < 10^{-4}$) but not at $q = 0.05$ ($p = 0.088$), so under the pre-registered rule these two predictions were evaluated at the two faster rates.

All four predictions were supported (Table 2, Fig 4). With an informative cue, fluctuation increased total regulatory strength relative to the frozen environment 1.8-fold at $q = 0.10$ and 2.7-fold at $q = 0.20$ (P1; Holm-adjusted $p = 0.005$ and $p \approx 10^{-4}$). Without a cue, fluctuation lowered regulatory strength about fivefold at all three rates (1.1 to 1.2 compared with 6.3; P3; Holm-adjusted $p \approx 10^{-4}$ at each rate). The informative-cue cells at the slowest rate, which were excluded from the verdict, showed the same direction (11.0 compared with 5.9). In one of those populations every regulatory connection was lost for the whole analysis window, and no other population in the sweep lost them completely.

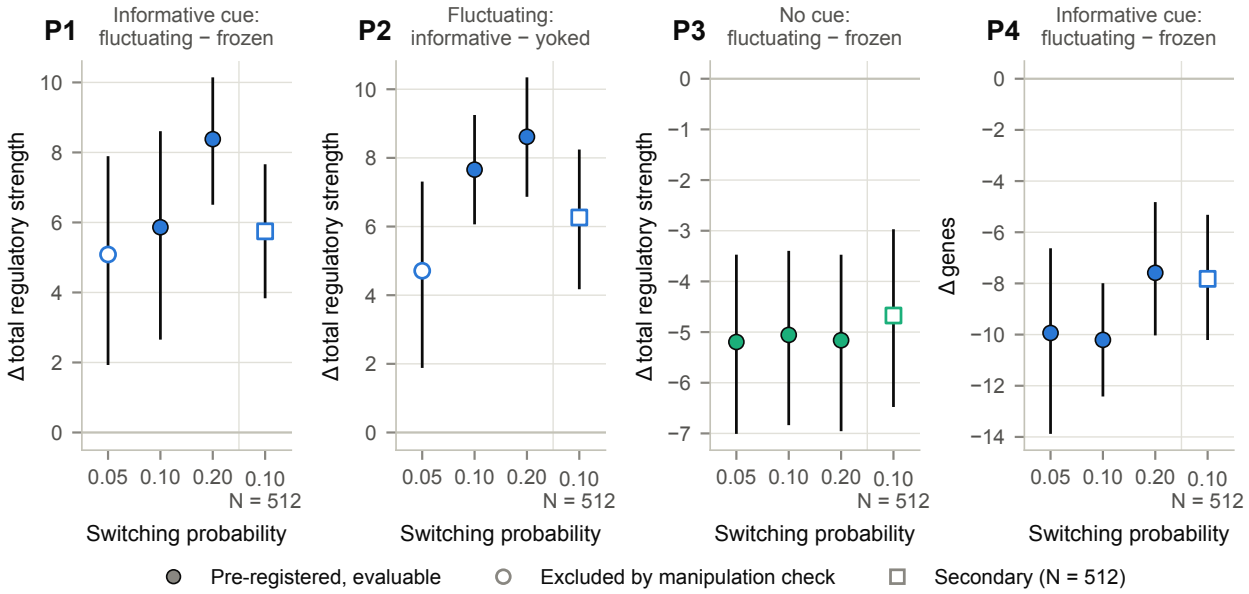



**Fig 4. Pre-registered test.** Differences in means with bootstrap 95% confidence intervals for each pre-registered contrast, at each switching probability (population 1024) and at population 512 (squares; secondary, no verdict). Open circles mark cells excluded from the verdict because the manipulation check failed at that rate. For P3 at population 512 the frozen baseline is the population-1024 frozen arm, the only one run. P1: effect of fluctuation on total regulatory strength with an informative cue. P2: informative minus yoked cue under fluctuation. P3: effect of fluctuation without a cue. P4: effect of fluctuation on the number of functional genes.

## The information carried by the cue produces the effect

The yoked control separates information from stimulation. Informative and yoked arms receive the same two signal proteins at the same concentration and frequency, with the same per-age rates of target change and signal change, and differ only in whether the signal predicts the target. Informative-cue populations evolved 2.3 times ($q = 0.10$) and 2.9 times ($q = 0.20$) the regulatory strength of yoked-cue populations (P2; Holm-adjusted $p \approx 10^{-4}$ in each case) and matched the target about twice as closely (metabolic error 0.029 against 0.058 at $q = 0.10$). They also carried more transcription factors at every rate (11.2 against 6.2 in the exploratory experiment and 10.6 against 7.4 at $q = 0.10$ in the sweep). In the frozen environment the signal still switches but predicts nothing, because the target does not change. There the informative and

yoked cues did not differ significantly in the pre-registered sweep (7.5 and 5.4, $p = 0.18$), and in the exploratory experiment the difference went the other way (4.7 and 6.2). The difference between cue arms appears only when the cue carries information about a changing environment.

The information did not increase genome size. Informative- and yoked-cue populations carried similar numbers of functional genes in the exploratory experiment (19.0 and 17.9, $p = 0.30$) and in the sweep (16.3 and 16.4, 18.1 and 19.9, and 18.0 and 20.8 at the three rates), where the only significant difference, at $q = 0.20$, favoured the yoked cue ($p = 0.01$, uncorrected). This agrees with the earlier EvoSim conclusion, which held even in a regime in which organisms read the cue. With reversible development, informative-cue populations matched the best-paying resource channel in 1.5 of 6 states, against 1.0 with a scrambled cue, yet had the same number of nodes (S4 Text). What information changed in R-Aevol was regulation, which has not been measured in that EvoSim regime.

## Fluctuation reduces genome size

With an informative cue, fluctuation reduced the number of functional genes at every switching rate, by 10, 10 and 8 genes (26.3, 28.3 and 25.6 in frozen populations compared with 16.3, 18.1 and 18.0 in fluctuating ones; P4; Holm-adjusted $p \leq 3.2 \times 10^{-4}$; Fig 4). The reduction was also present at a population size of 512 (25.3 compared with 17.5, $p < 10^{-4}$), and without a cue or with the yoked cue (S3 Table). Other models have reported the opposite, larger genomes under fluctuating selection [12,13], and so did earlier EvoSim experiments (S4 Text). A possible explanation lies in how Aevol builds phenotypes. The phenotype is the sum of contributions from individual genes, so a constant target that can be matched precisely favours additional genes. A target that changes within a lifetime cannot be matched precisely, and the coarser phenotype it leaves needs fewer genes. Consistent with this, frozen populations reached a metabolic error of about 0.008 and fluctuating populations without a usable cue about 0.06. We have not tested this explanation directly, and two results do not support it. Coarse-grained populations matched their targets better than fine-grained ones (metabolic error 0.038 against 0.062) but did not carry more genes, and informative-cue populations matched theirs better than cue-free ones (0.029 against 0.063 at $q = 0.10$) with fewer genes (18.1 against 20.0).

## Connection counts do not track connection strength

Connection number and connection strength did not always change together (Fig 5). In EvoSim, fluctuation increased the number of connections by 39% and left total wiring strength unchanged. In R-Aevol, fluctuating populations with an informative cue had no more regulatory links than the corresponding frozen populations (0.60, 0.78 and 0.83 times as many at the three rates; $p = 0.02$, 0.13 and 0.52), although their total regulatory strength was 1.8 to 2.7 times higher and their strength per link about twice as high. A study measuring connectivity as link number would have missed the strengthening of regulation, and its point estimates would have suggested the opposite. Without a cue, number and strength fell together to about a fifth of the frozen values, so the two measures do not always disagree. These comparisons were pre-specified as secondary analyses without a formal verdict.

## Environmental grain

Tsuda and Kawata [12] found that fluctuation across generations, with no information about the position of the optimum, promotes regulatory complexity. The effect was strong when the optimum moved in a random walk, and slight when it alternated between two positions as in our design. Their fluctuation was also slow relative to a lifetime, whereas ours was fast. To test whether grain explains the difference, we varied it in R-Aevol without modifying the model. The environment is a single Markov chain over ages that continues from one generation to the next, so the switching probability sets how long an environment persists on average, from half a lifetime at $q = 0.1$ to about 250 generations at $q = 0.0002$ (Fig 1D). We predicted that coarse-grained cue-free fluctuation would produce larger genomes and stronger regulation than fine-grained fluctuation (G1, G3) and than a frozen environment (G2, G4), the direction reported by

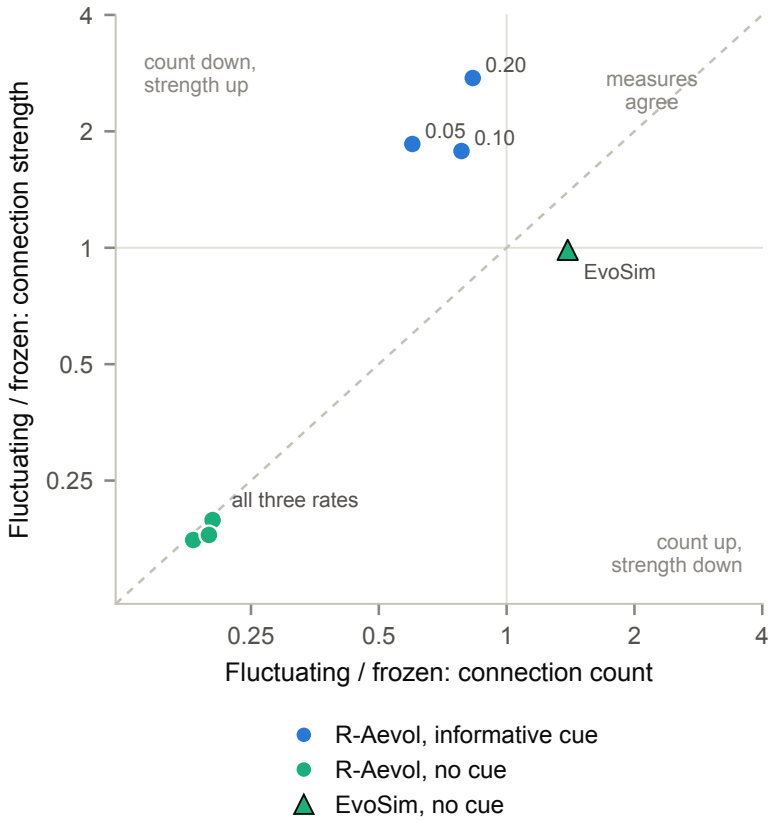



**Fig 5. Counting versus weighing.** Ratio of fluctuating to frozen connection count against the corresponding ratio of connection strength, on logarithmic axes. R-Aevol (circles): raw number of regulatory links and total regulatory strength, per switching probability, population 1024. EvoSim (triangle): number of non-zero connections and total wiring strength, no cue. Points are ratios of means. Points on the dashed diagonal are cases in which the two measures agree, and points off it are cases in which counting misreports strength. For R-Aevol with an informative cue, the difference in link number is significant only at $q = 0.05$.

Tsuda and Kawata. How each possible outcome would be interpreted was fixed in advance (Materials and methods).

The results did not support this prediction (Table 2, Fig 6). Under cue-free fluctuation, genome size and regulatory strength showed no detectable dependence on grain. At the coarsest grain, populations carried 19.2 genes and had a regulatory strength of 2.0, against 19.5 and 2.3 at the finest (G1 and G3, inconclusive; $p = 0.80$ and 0.60), and a secondary analysis found no trend across the four grains (Spearman correlation with log dwell time 0.06 for genes and 0.01 for strength). Coarse-grained fluctuation reduced gene number and regulatory strength relative to the frozen environment by about as much as fine-grained fluctuation did (25.6 genes and a strength of 8.4 in frozen populations; G2 and G4 refuted, Holm-adjusted $p = 4 \times 10^{-4}$ and $8 \times 10^{-5}$). Coarser grain did improve the match to the target, lowering metabolic error from 0.062 to 0.038, but genome size and regulation stayed where they were. This is the pattern expected of adaptive tracking, in which populations re-adapt genetically to each environment when it lasts many generations [9, 27]. The finest-grained cell, with new seeds, reproduced P3 and the reduction in gene number without a cue. By the interpretation fixed in advance, grain does not reconcile our results with those of Tsuda and Kawata.

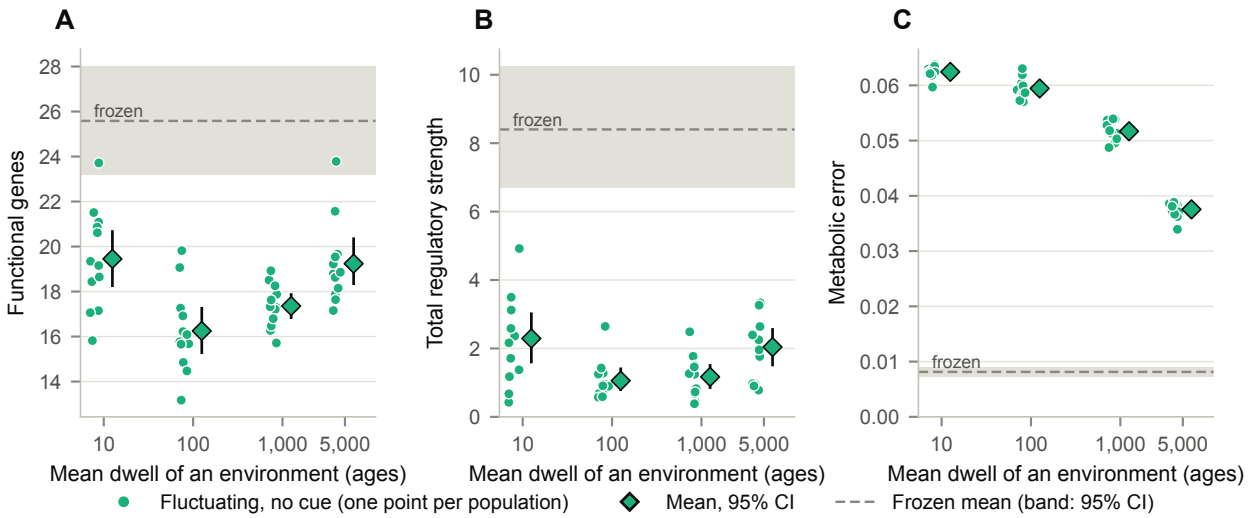



**Fig 6. Environmental grain.** Cue-free fluctuating populations at four switching probabilities, plotted against the mean time an environment persists ($1/q$ ages; 20 ages make one generation), one point per population, with means and bootstrap 95% confidence intervals. The dashed line and grey band show the mean and 95% confidence interval of the frozen arm. (A) Functional genes. (B) Total regulatory strength. (C) Metabolic error, the distance between phenotype and target.

## Discussion

In R-Aevol, the effect of environmental fluctuation on regulation depended on whether organisms could identify the current environment. Fluctuation with an informative cue produced the strongest regulation we observed and fluctuation without a cue the weakest. A yoked cue, which supplied the same molecules at the same rates, did not have the effect of an informative one. These results held across switching rates and at two population sizes, and the cost of cue-free fluctuation held across a 500-fold range of environmental grain, in tests whose predictions and analyses were fixed before the data were generated.

The phenotypic side of this result is well known. Reliable cues favour plastic responses, and in their absence generalist phenotypes or bet-hedging are favoured [6–9]. This has recently been shown for the evolution of plasticity in a gene regulatory network, in which the parts of the network that become sensitive to mutation also depend on the environmental regime [11], and in a network with a dedicated sensor gene, plasticity evolves as the sensor comes to activate some genes and repress others [25]. Our results add the effect of a cue's information on the regulatory strength and the size of the network that implements the response. The yoked control separated that information from the cue's physical presence, which a manipulation of cue reliability achieves only if it includes a cue that is present at the usual frequencies but uninformative.

We designed the grain experiment to test whether grain accounts for the difference from Tsuda and Kawata, who found that fluctuation of the optimum across generations promotes regulatory complexity [12]. It does not. In R-Aevol, cue-free fluctuation reduced genome size and regulatory strength by similar amounts whether an environment lasted half a lifetime or about 250 generations, and coarser grain improved how closely populations matched their targets without detectably changing either quantity. Under the interpretation we fixed in advance, the remaining difference between the two studies is therefore one between the models. Going beyond that interpretation, Tsuda and Kawata's own results point to the pattern of fluctuation, which we had listed as a limitation of the design when registering it. Their strong effect came from an optimum that moved to new positions in a random walk. In their cyclic regime, which alternated between two optima as ours does, complexity rose only slightly and only for large, infrequent shifts [12], so for this kind of fluctuation the two studies differ less than their headline results suggest. A design with novel optima would test this directly. The way each model maps genotype to phenotype is a further candidate. The same applies to EvoSim, which produced larger genomes under fluctuation in

earlier experiments. Its fluctuation is coarse relative to a lifetime, since development happens once at birth and each environmental state lasts about 14 generations, but coarse grain did not enlarge genomes in R-Aevol, so grain is an unlikely explanation, and the difference in genome size between EvoSim and R-Aevol is most likely a difference between the models.

The divergence between connection number and connection strength probably has different causes in the two models, and we have tested neither. In EvoSim, gene duplication and node addition copy the connections of an existing node and add small random values to every copied weight, including weights that were zero, so each duplication can create many weak connections. Point mutations that hit a zero weight also create connections. If purifying selection is weaker in the smaller populations of fluctuating environments, such connections would persist and connection number would rise without an increase in total strength, which fits the correlations with population size within the fluctuating arms. In R-Aevol, frozen populations had at least as many regulatory links as informative-cue populations, but weaker ones, and we did not determine why. Many evolutionary models of regulation measure complexity by counting genes or connections [12, 17], and we recommend that weighted measures be reported as well.

## Limitations

The positive result, that a usable cue allows fluctuation to strengthen regulation, rests on R-Aevol alone. EvoSim organisms respond to a cue only when development is reversible and even then only in part, and we have not measured regulatory strength in that regime. Each R-Aevol condition contained 12 populations, fewer than the 60 to 100 used in comparable studies [12, 28]. The effects were large ($|d|$ between 1.5 and 3.9 in the evaluated cells) and remained significant after correction for multiple testing, but smaller effects could have been missed. Because R-Aevol counts the influences of signal proteins as regulatory links, comparisons between arms with and without a cue include links that the no-cue arms cannot have. The pre-registered contrasts compare arms with the same signals and are not affected. Frozen populations without signals reached a regulatory strength of 6.3, within the range of frozen populations with signals (4.8 to 7.5), so regulation of that strength does not require signals. We cannot tell how much of the difference between fluctuating arms with and without a cue the signal links account for. The version of R-Aevol we used records per-generation statistics only for the best individual in each population, so every R-Aevol measure describes the fittest individual of each generation, averaged over the analysis window, and not the population average. R-Aevol's fluctuating environments alternate between two fixed targets and do not include the random walk towards new optima that Tsuda and Kawata found most favourable to complexity [12], and even at the coarsest grain a switch falls within about 0.4% of lifetimes. Prediction P4 establishes the direction of the effect on genome size; its mechanism is untested.

In EvoSim, fluctuation and population size are confounded between arms. The within-arm correlations argue against population size as the cause of the lower strength density but are consistent with a contribution to the higher connection count. We also pre-registered a confirmation of the EvoSim divergence between count and strength in a regime with more hidden nodes, on condition that fluctuating populations survived there. They did not (2 to 5 of 6 survived, compared with 6 of 6 frozen populations), so the confirmation was not run and the EvoSim result remains exploratory. Finally, all pre-registrations were committed to a version-control repository before the corresponding runs, but these timestamps exist only locally and cannot be verified independently.

## Conclusion

Environmental fluctuation did not by itself strengthen gene regulation in these models. Regulation became stronger under fluctuation when a cue allowed organisms to track the environment. Without a cue it became weaker, whatever the grain of the environment. Because connection counts and connection strength can change independently, conclusions about regulatory complexity may depend on which of the two is measured.

# Materials and methods

## EvoSim

EvoSim is described in full, following the ODD protocol [29], in S1 Text, and the parts relevant to this study are summarised below. Organisms live on a $G \times G$ lattice ($G = 128$) with six diffusing resource channels. External inflow, which varies seasonally, enters the first channel (with the exception described under Environment and cue), and consuming channel $r$ releases a byproduct into channel $r + 1$, with yields declining along this chain. All arms also experience rare random mass-mortality and famine events, on average one every 25,000 time steps. Each organism has an energy store and divides when the store exceeds a heritable threshold. Starvation and the other causes of death are described in S1 Text.

**Genome and development.** A genome is a weight matrix $W$ and a bias vector $b$ over $n_g$ nodes, with $13 \le n_g \le 32$. Nodes 1 to 12 are read out as the phenotype, node 13 receives the abiotic cue, and hidden nodes $14, \ldots, n_g$ have no fixed meaning. Development takes place once, at birth. Expression is initialised at 0.1 and updated eight times as

$$e_k \leftarrow \tanh\Big(b_k + \sum_m W_{mk}\, e_m + u_k\Big), \tag{1}$$

where $u_k$ is the local concentration of resource $k$ for $k \le 6$, the cue for $k = 13$, and zero otherwise. The final expression values are mapped to the twelve traits.

**Mutation.** At each birth, a Poisson number of point mutations with mean $10^{-3} n_g^2$ is applied, each adding $\mathcal{N}(0, 0.25^2)$ to a randomly chosen weight. This rate is reduced by up to 85% by the organism's heritable repair trait. With probability equal to the reduced per-weight rate times $n_g$, one node bias receives a mutation of the same size. The structural mutations are gene silencing (probability 0.002), which sets a node's outgoing weights to zero; gene duplication (0.002), which copies the regulatory role of one node onto another and adds $\mathcal{N}(0, 0.05^2)$ noise to every copied weight; and the addition or deletion of a hidden node (0.003 each), where an added node is a duplicate of an existing one made in the same way. Connection silencing (probability 0.2 per birth) sets one randomly chosen weight to exactly zero. Without it, additive mutation would drive density towards one. Each hidden node costs 0.0015 and each non-zero weight $5 \times 10^{-5}$ energy units per time step.

**Environment and cue.** An abiotic condition cycles through discrete states with a period of 2,000 time steps. Each state changes which resource channel gives the highest yield, and in fluctuating arms 40% of the external inflow goes to the favoured channel. In frozen arms all inflow enters channel 1, which then also gives the highest yield. The phase of the cycle, scaled to $[-1, 1]$, is the input to node 13 and identifies the current state. Fluctuating arms have six states and frozen arms one. In the yoked arm the cue is instead the phase multiplied by $10/\varphi \approx 6.18$, with $\varphi$ the golden ratio, and wrapped back into the same range, so that it takes the same values but carries only 0.0004 bits of information about the state. In the no-cue arm the gain of the cue is zero. With a single state the cue is identically zero, so the three frozen arms produce identical runs and were analysed as one arm.

**Measures.** For every organism at birth we record the number of weights with $|w| > 10^{-4}$, the strength density $\sum |w| / n_g^2$ over all $n_g^2$ weights, including zeros, and $n_g$. The log holds their population means at each time step. Count density is the mean number of non-zero weights divided by the squared mean $n_g$, and total strength is the mean strength density multiplied by the squared mean $n_g$. Each measure is then averaged over the last 20% of the run. Populations that went extinct were excluded and their number is reported.

**Runs.** We ran six arms with 30 seeds each for 150,000 time steps (180 runs), about 6,200 generations in fluctuating and 4,900 in frozen populations (24 and 31 time steps per generation, from the slope of mean lineage depth). Runs are deterministic for a given seed and, on the same machine and Julia version, bit-identical across thread counts (checked at 1, 4 and 8 threads).

## R-Aevol

We used the `aevol-7` branch of Aevol (commit `33277374`, https://gitlab.inria.fr/aevol/aevol) compiled with regulation enabled (`-Dwith-raevol=ON`) and with per-generation statistics (`-Dwithout-stats=OFF`). Building on Apple Silicon required correcting a hard-coded path to the OpenMP library (S2 Text). The parameters (Table 1) are taken from the example configurations distributed with the platform (`basic_regul`, `dynamic_env` and `complex_env`). The mutation and rearrangement rates are those of `complex_env`, ten times those of `dynamic_env`, and the binding matrix is the file supplied with `dynamic_env`, used unchanged in every run.

**Environments.** Each environment is a target function made of three Gaussians. Target B is target A reflected about the midpoint of the phenotypic axis, and the two have the same geometric area, 0.152468. Fluctuating arms switch between A and B. Frozen arms use two copies of the same target (A for odd seeds, B for even seeds), so that the switching mechanism runs in the same way but the organism never experiences a change. The environment is a Markov chain over ages that continues across generations. At each of the 20 ages it moves, with probability $q$, to a different environment chosen uniformly among the others. We checked in the source code that this is the switching carried out by the platform's SIMD engine (S2 Text).

**Signals and the yoked control.** Two signal proteins, taken unchanged from the `complex_env` example, are present at concentration 0.9 during the ages in which their environment is active and absent otherwise. The informative arm has two environments, (A, s1) and (B, s2). The yoked arm has four, (A, s1), (A, s2), (B, s1) and (B, s2). The no-cue arm has no signals. In the two-environment arm every switch changes the target and the signal together, so each changes at rate $q$ per age. In the four-environment arm a switch changes the target with probability 2/3 and the signal with probability 2/3, so with a switching probability of $1.5q$ each again changes at rate $q$ per age.

**Measures.** R-Aevol records statistics of the best individual in each generation. Genome size is the number of functional genes. Total regulatory strength is $S = n_{\text{link}}\,\bar{a}$, where $n_{\text{link}}$ is the number of activating and repressing protein–RNA influences (a protein that activates and also represses an RNA counts twice) and $\bar{a}$ is their mean strength (the sum of the activation and repression factors divided by their number). $S = 0$ when there are no links. R-Aevol treats the signal proteins like other proteins when it computes binding, so in arms with a cue $n_{\text{link}}$ and $S$ include influences of the signal proteins, which the statistics do not separate. Each run is summarised by its mean over generations 25,001 to 30,000.

**Runs.** The exploratory experiment had six arms with 12 seeds each at $q = 0.10$ and a population of 1024 (72 runs). The pre-registered sweep comprised fluctuating informative, yoked and no-cue arms and frozen informative arms at $q = 0.05$, 0.10 and 0.20, together with frozen no-cue and frozen yoked arms at $q = 0.10$ (the frozen no-cue arm does not depend on $q$), each with 12 new seeds and a population of 1024 (168 runs). The fluctuating informative, yoked and no-cue arms and the frozen informative arm were repeated at a population of 512 (48 runs). The grain experiment has cue-free fluctuating arms at $q = 0.1$, 0.01, 0.001 and 0.0002 and a frozen arm, with 12 new seeds each (60 runs). All runs reached 30,000 generations.

## Statistics

The unit of analysis is a population. Each run contributes one value per measure, so generations within a run are not treated as independent. Differences between arms were tested with two-sided permutation tests on the difference of means (100,000 permutations). Rank-based permutation tests of the pre-registered contrasts led to the same verdicts (S2 Text). Multiple predictions were corrected with the Holm procedure. Effect sizes are given as Cohen's $d$, and intervals in the figures are bootstrap 95% confidence intervals (10,000 resamples).

**Table 1. R-Aevol parameters.**

| Parameter | Value |
|---|---|
| Population | 1024 ($32 \times 32$) or 512 ($32 \times 16$) |
| Selection | fitness-proportionate, pressure 750 |
| Point mutation, small insertion, small deletion | $5 \times 10^{-5}$ per base per generation |
| Duplication, deletion, translocation, inversion | $5 \times 10^{-5}$ per base per generation |
| Maximum indel size | 6 |
| Initial genome | 5,000 bases, one functional gene |
| Hill function | $n = 4$, $\theta = 0.5$ |
| Degradation | rate 1, 10 steps per age |
| Lifetime | 20 ages, fitness evaluated at all 20 |
| Binding matrix | fixed file from `examples/dynamic_env`, all runs |
| Target sampling | 300 points |
| Signal concentration | 0.9 |
| Generations | 30,000 |

## Pre-registration and deviations

Predictions P1 to P4, their measures, tests and decision rules, and the analysis code were committed to version control (commit `f985ccd`) before the pre-registered sweep was started. The analysis was run once, on the complete data, with the committed code, whose checksum was verified beforehand. The only deviation was operational. The idle-sleep setting of the host machine suspended the runs overnight, which lengthened the wall time but cannot have affected the results, because a suspended process resumes from the state in which it stopped. The EvoSim confirmation (commit `4bc16c3`) required fluctuating and frozen populations to survive a viability pilot. None of the three candidate regimes met this condition, so it was not run (commit `56fe3a0`). The design, predictions and analysis code of the grain experiment were committed in the same way (commit `fa00777`) before its runs began, and its analysis was run once, automatically, when the last run finished. Its interpretation was fixed at the same time. If coarse grain gave more genes (or stronger regulation) than fine grain and than the frozen environment, grain would reconcile the two studies. If it did so only relative to fine grain, Tsuda and Kawata's gain would probably need novel optima, which our design cannot produce. If it did not do so relative to fine grain, grain would not be the explanation and the difference would be reported as one between the models.

## Software and availability

EvoSim is written in Julia and uses only the standard library for simulation. The R-Aevol runs were made on a single 12-core machine. The configurations of the two models, the launch scripts, the pre-registered analysis scripts, the history of all analyses and the per-run statistics are available at **[PENDING:** repository URL and Zenodo DOI].

# Supporting information

**S1 Text. ODD description of EvoSim.** Complete model description following the ODD protocol, with every parameter value.

**S2 Text. R-Aevol build, configuration and the yoked-control derivation.** Build options and the changes needed on Apple Silicon, the parameter-file template, targets and signals, the environment-switching code, the derivation of matched rates of target and signal change, the definition of the measures, seeds, the pre-registration record, and the result of every pre-registered test with rank-based $p$ values.

**Table 2. Pre-registered predictions and verdicts.**

| Id | Prediction | Verdict |
|---|---|---|
| P1 | Fluctuation with an informative cue raises total regulatory strength over frozen | supported ($q = 0.10$, 0.20) |
| P2 | An informative cue raises it over a yoked cue | supported ($q = 0.10$, 0.20) |
| P3 | Fluctuation without a cue lowers it below frozen | supported (all rates) |
| P4 | Fluctuation with an informative cue lowers the number of functional genes | supported (all rates) |
| P5 | In EvoSim, count density rises and strength density falls in a regime above the node floor | not run (no viable regime) |
| G1, G3 | Coarse-grained cue-free fluctuation gives more genes (G1) and stronger regulation (G3) than fine-grained | inconclusive |
| G2, G4 | Coarse-grained cue-free fluctuation gives more genes (G2) and stronger regulation (G4) than frozen | refuted |

A prediction was supported if the predicted sign held at every evaluable rate and the Holm-adjusted $p < 0.05$ at two or more, refuted if the opposite sign was significant at any rate or no rate was Holm-significant, and inconclusive otherwise. P1 and P2 are evaluable only at rates where the informative cue measurably lowers metabolic error relative to the yoked cue. G1–G4 are single tests, supported if the predicted sign holds with Holm-adjusted $p < 0.05$, refuted if the opposite sign holds with $p < 0.05$, and inconclusive otherwise.

**S3 Table. Results for every arm of every experiment.** Means and standard deviations for all 528 runs. Per-run values are in the repository.

**S4 Text. Earlier EvoSim results and the audit trail.** Status of earlier claims after this study, the record of pre-registered predictions, and the measurement errors found during the project with the checks that exposed them.

**S5 Table. Prior work and how this study relates to it.**

## Acknowledgments

[**PENDING:** acknowledgements; disclosure of AI-assisted analysis and writing, as required by PLOS policy]

## References

1. Konstantinidis KT, Tiedje JM. Trends between gene content and genome size in prokaryotic species with larger genomes. Proc Natl Acad Sci USA. 2004;101(9):3160-5.
2. van Nimwegen E. Scaling laws in the functional content of genomes. Trends Genet. 2003;19(9):479-84.
3. Cases I, de Lorenzo V, Ouzounis CA. Transcription regulation and environmental adaptation in bacteria. Trends Microbiol. 2003;11(6):248-53.
4. Galperin MY. A census of membrane-bound and intracellular signal transduction proteins in bacteria: bacterial IQ, extroverts and introverts. BMC Microbiol. 2005;5:35.
5. Beslon G, Parsons DP, Sanchez-Dehesa Y, Peña JM, Knibbe C. Scaling laws in bacterial genomes: a side-effect of selection of mutational robustness? BioSystems. 2010;102(1):32-40.
6. Levins R. Evolution in Changing Environments. Princeton: Princeton University Press; 1968.

7. Moran NA. The evolutionary maintenance of alternative phenotypes. Am Nat. 1992;139(5):971-89.

8. Kussell E, Leibler S. Phenotypic diversity, population growth, and information in fluctuating environments. Science. 2005;309(5743):2075-8.

9. Botero CA, Weissing FJ, Wright J, Rubenstein DR. Evolutionary tipping points in the capacity to adapt to environmental change. Proc Natl Acad Sci USA. 2015;112(1):184-9.

10. Ashander J, Chevin LM, Baskett ML. Predicting evolutionary rescue via evolving plasticity in stochastic environments. Proc R Soc B. 2016;283(1839):20161690.

11. Jorritsma A, van den Berg P. The evolution of plasticity and evolvability in a simple gene regulatory network. Evolution. 2026;80(1):230-9. doi:10.1093/evolut/qpaf219.

12. Tsuda ME, Kawata M. Evolution of gene regulatory networks by fluctuating selection and intrinsic constraints. PLoS Comput Biol. 2010;6(8):e1000873.

13. Bentkowski P, van Oosterhout C, Mock T. A model of genome size evolution for prokaryotes in stable and fluctuating environments. Genome Biol Evol. 2015;7(8):2344-51.

14. Malcom JW. Smaller gene networks permit longer persistence in fast-changing environments. PLoS One. 2011;6(4):e14747.

15. Knibbe C, Coulon A, Mazet O, Fayard JM, Beslon G. A long-term evolutionary pressure on the amount of noncoding DNA. Mol Biol Evol. 2007;24(10):2344-53.

16. Luiselli J, Rouzaud-Cornabas J, Lartillot N, Beslon G. Genome streamlining: effect of mutation rate and population size on genome size reduction. Genome Biol Evol. 2024;16(12):evae250.

17. Espinosa-Soto C. On the role of sparseness in the evolution of modularity in gene regulatory networks. PLoS Comput Biol. 2018;14(5):e1006172.

18. Kashtan N, Alon U. Spontaneous evolution of modularity and network motifs. Proc Natl Acad Sci USA. 2005;102(39):13773-8.

19. Ng ETH, Kinjo AR. Plasticity-led evolution as an intrinsic property of developmental gene regulatory networks. Sci Rep. 2023;13(1):19830.

20. Ng ETH, Kinjo AR. Plasticity-led and mutation-led evolutions are different modes of the same developmental gene regulatory network. PeerJ. 2024;12:e17102.

21. Church RM. Systematic effect of random error in the yoked control design. Psychol Bull. 1964;62:122-31.

22. Soyer OS, Bonhoeffer S. Evolution of complexity in signaling pathways. Proc Natl Acad Sci USA. 2006;103(44):16337-42.

23. Iwasaki WM, Tsuda ME, Kawata M. Genetic and environmental factors affecting cryptic variations in gene regulatory networks. BMC Evol Biol. 2013;13:91.

24. Sanchez-Dehesa Y, Parsons D, Peña JM, Beslon G. Modelling evolution of regulatory networks in artificial bacteria. Math Model Nat Phenom. 2008;3(2):27-66.

25. Sakamoto T, Innan H. On the evolutionary origin of discrete phenotypic plasticity. G3 (Bethesda). 2024;14(9):jkae144.

26. Luiselli J, Parsons DP, Gallé R, Banse P, Rouzaud-Cornabas J, Beslon G. Aevol-9: a simulation platform to decipher the evolution of genome architecture. bioRxiv [Preprint]. 2025. doi:10.1101/2025.04.10.648095.

27. Rago A, Kouvaris K, Uller T, Watson R. How adaptive plasticity evolves when selected against. PLoS Comput Biol. 2019;15(3):e1006260.

28. Lalejini A, Ferguson AJ, Grant NA, Ofria C. Adaptive phenotypic plasticity stabilizes evolution in fluctuating environments. Front Ecol Evol. 2021;9:715381.

29. Grimm V, Railsback SF, Vincenot CE, Berger U, Gallagher C, DeAngelis DL, et al. The ODD protocol for describing agent-based and other simulation models: a second update to improve clarity, replication, and structural realism. J Artif Soc Soc Simul. 2020;23(2):7.